

In [5]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [6]:

```
df=pd.read_csv(r"C:\Users\91949\Downloads\fiat500_VehicleSelection_Dataset.csv")
df
```

Out[6]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.496
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns

In [7]:

```
df=df[['engine_power', 'age_in_days']]
df.columns=['ep', 'aid']
```

In [8]:

```
df.describe()
```

Out[8]:

	ep	aid
count	1538.000000	1538.000000
mean	51.904421	1650.980494
std	3.988023	1289.522278
min	51.000000	366.000000
25%	51.000000	670.000000
50%	51.000000	1035.000000
75%	51.000000	2616.000000
max	77.000000	4658.000000

In [9]:

```
df.head()
```

Out[9]:

	ep	aid
0	51	882
1	51	1186
2	74	4658
3	51	2739
4	73	3074

In [10]:

```
df.tail()
```

Out[10]:

	ep	aid
1533	51	3712
1534	74	3835
1535	51	2223
1536	51	2557
1537	51	1766

In [11]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    ep      1538 non-null    int64  
 1    aid      1538 non-null    int64  
dtypes: int64(2)
memory usage: 24.2 KB
```

In [12]:

```
df.isna().any()
```

Out[12]:

```
ep      False
aid      False
dtype: bool
```

In [13]:

```
df.fillna(method='ffill',inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_11540\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method='ffill',inplace=True)
```

In [14]:

```
x=np.array(df['ep']).reshape(-1,1)
y=np.array(df['aid']).reshape(-1,1)
```

In [15]:

```
df.dropna(inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_11540\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
df.dropna(inplace=True)
```

In [16]:

```
sns.lmplot(x="ep",y="aid",data= df,order=2,ci=None)
```

Out[16]:

<seaborn.axisgrid.FacetGrid at 0x1ddce7c2b90>




In [17]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```

0.09045713743251182

In [18]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```




In [20]:

```
df100=df[:][:100]  
sns.lmplot(x='ep',y='aid',data=df100,order=1,ci=None)
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x1ddd50c2bd0>




In [21]:

```
df100.fillna(method='ffill',inplace=True)
X=np.array(df100['ep']).reshape(-1,1)
y=np.array(df100['aid']).reshape(-1,1)
df100.dropna(inplace=True)
X_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr=LinearRegression()
regr.fit(X_train,y_train)
print(regr.score(x_test,y_test))
print("Regression: ",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

0.10340082129282724

Regression: 0.10340082129282724



In [22]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(X_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2_score: ",r2)
```

R2_score: 0.10340082129282724

Conclusion:

Dataset we have taken is poor for linear model but with the smaller data works well linear model